

# PROBLEM SET 4 – NEURAL NETWORKS

Ravi Teja Sunkara (rsunkara) (50292191)  
UNIVERSITY AT BUFFALO (rsunkara@buffalo.edu)

## Question 1

Ans Neural networks with at least 1 hidden layer are universal approximators, which means that they can approximate any (continuous) function. This approximation can be improved by increasing the number of hidden neurons in the network (but increases the risk of overfitting).

- A key advantage to neural networks is that they are capable of learning features independently, with much less human involvement.

Softmax function :-

The softmax function (called such because it is like a "softened" maximum function) may be used as the output layer's activation function. It takes the form:

Softmax is usually used for multivariate logistic regression because it produces a categorical distribution by squashing activation values to be between 0 & 1 and sum to 1. We have used it to implement a different type of plenty (entropy-based) on distributions.

- This function has the properties that it sums to 1 and that all of its outputs are +ve, which are useful for modeling probability distributions.

- The cost function to use with softmax is the (categorical) cross-entropy loss function. It has the nice property of having a very high gradient when the target value is 1 and the output is almost 0.

## Negative Log-Likelihood:-

In practice, the softmax function is used in tandem with the -ve log-likelihood. This loss function is very interesting if we interpret it in relation to the behavior of softmax.

First, let's write down our loss functions:

This is summed for all the correct classes.

- Recall, that when training a model, we expect to find the minima of a loss function given a set of parameters (in a neural network, these are weights and biases). We can interpret the loss as the "unhappiness" of the network with respect to its parameters. The higher the loss, the higher the unhappiness; we don't want that. We want to make our model happy.

For example:

- Let us assume that we have  $N$  images and  $y_i$  is the ~~image of the~~ label of the image  $i$ , where  $y_i$  contains  $R \times L$  - a binary vector of length  $C$  (no. of classes)  $y_i \in L$  when the image is belonging to class  $C$ .

Consider the following two log functions

$$L1 = - \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(P(y_{i,c}|D))$$

$$L2 = \sum_{i=1}^N \sum_{c=1}^C (y_{i,c} - P(y_{i,c}|D))^2$$

- $L2$  is not always used with neural networks, indeed for statistical pattern recognition problems the cross-entropy loss (with a softmax activation function for the output layer) is the preferred option.

Ans Neural networks with at least 1 hidden layer are universal approximators, which means that they can approximate any (continuous) function. This approximation can be improved by increasing the number of hidden neurons in the network (but increases the risk of overfitting).

→ A key advantage to neural networks is that they are capable of learning features independently, with much human involvement.

**Softmax function :-**  
 The softmax function (called such because it is like a "softened" maximum function) may be used as the output layer's activation function. It takes the form:  
 Softmax is usually used for multivariate logistic regression because it produces a categorical distribution by squashing activation values to be between 0 & 1 and sum to 1. We have used it to implement a different type of plenty (entropy-based) on distributions.  
 → This function has the properties that it sums to 1 and that all of its outputs are rve, which are useful for modeling probability distributions.  
 → The cost function to use with softmax is the (categorical) cross-entropy loss function. It has the nice property of having a very big gradient when the target value is 1 and the output is almost 0.

Other than the provided answer the below claims also support the argument:

The negative log likelihood (eq.80) is also known as the multiclass cross-entropy (ref: Pattern Recognition and Machine Learning Section 4.3.4), as they are in fact two different interpretations of the same formula.

eq.57 is the negative log likelihood of the Bernoulli distribution, whereas eq.80 is the negative log likelihood of the multinomial distribution with one observation (a multiclass version of Bernoulli).

For binary classification problems, the softmax function outputs two values (between 0 and 1 and sum to 1) to give the prediction of each class. While the sigmoid function outputs one value (between 0 and 1) to give the prediction of one class (so the other class is 1-p).

So eq.80 can't be directly applied to the sigmoid output, though it is essentially the same loss as eq.57.

## Question 2 (a)

Build a neural network with 1 hidden layer of 30 sigmoid nodes, and an output layer 10 softmax nodes from 1000 training images (100 images per digit). Train the network for 30 complete

epochs, using mini-batches of 10 training examples at a time, a learning rate  $\eta=0.1$ . Plot the training error, testing error, criterion function on training data set, criterion function on testing data set of a separate 1000 testing images (100 images per digit), and the learning speed of the hidden layer (the average absolute changes of weights divided by the values of the weights).

**Solution:**

The code for this is the following file:single\_hidden\_layer

Here are the specifications of this execution:

The following libraries have been used: tensorflow along with keras.

Also, the following requirements have been met in the code:

1. The code was run only on 1000 training and 1000 test images.
2. It had only one hidden layer.
3. The hidden layer had sigmoid as it's activation function.
4. There were 30 sigmoid nodes.
5. An output layer was there having 10 softmax nodes.
6. The network was trained for 30 complete epochs.
7. Also a mini batch of 10 training examples was used at a time
8. Learning rate was in initially kept as 0.1
9. In keras the learning rate speed can be changed using the decay parameter of the activation function

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 30)	23550
dense_4 (Dense)	(None, 10)	310

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

---

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 0s - loss: 2.2700 - acc: 0.1830 - val\_loss: 2.1821 - val\_acc: 0.4040

- LR: 0.090909

Epoch 2/30

- 0s - loss: 2.0282 - acc: 0.4160 - val\_loss: 1.8824 - val\_acc: 0.5900

- LR: 0.083333

Epoch 3/30

- 0s - loss: 1.6315 - acc: 0.6300 - val\_loss: 1.5300 - val\_acc: 0.5910

- LR: 0.076923

Epoch 4/30

- 0s - loss: 1.2944 - acc: 0.7160 - val\_loss: 1.2702 - val\_acc: 0.6580

- LR: 0.071429

Epoch 5/30

- 0s - loss: 1.0616 - acc: 0.7660 - val\_loss: 1.1003 - val\_acc: 0.7110

- LR: 0.066667

Epoch 6/30

- 0s - loss: 0.8979 - acc: 0.8050 - val\_loss: 0.9914 - val\_acc: 0.7340  
- LR: 0.062500

Epoch 7/30

- 0s - loss: 0.7893 - acc: 0.8270 - val\_loss: 0.8957 - val\_acc: 0.7540  
- LR: 0.058824

Epoch 8/30

- 0s - loss: 0.7050 - acc: 0.8470 - val\_loss: 0.8329 - val\_acc: 0.7620  
- LR: 0.055556

Epoch 9/30

- 0s - loss: 0.6441 - acc: 0.8520 - val\_loss: 0.7924 - val\_acc: 0.7790  
- LR: 0.052632

Epoch 10/30

- 0s - loss: 0.5927 - acc: 0.8660 - val\_loss: 0.7525 - val\_acc: 0.7810  
- LR: 0.050000

Epoch 11/30

- 0s - loss: 0.5545 - acc: 0.8730 - val\_loss: 0.7189 - val\_acc: 0.7940  
- LR: 0.047619

Epoch 12/30

- 0s - loss: 0.5210 - acc: 0.8760 - val\_loss: 0.6961 - val\_acc: 0.8000  
- LR: 0.045455

Epoch 13/30

- 0s - loss: 0.4930 - acc: 0.8840 - val\_loss: 0.6772 - val\_acc: 0.8000  
- LR: 0.043478

Epoch 14/30

- 0s - loss: 0.4687 - acc: 0.8900 - val\_loss: 0.6615 - val\_acc: 0.8010  
- LR: 0.041667

Epoch 15/30

- 0s - loss: 0.4477 - acc: 0.8970 - val\_loss: 0.6465 - val\_acc: 0.8100  
- LR: 0.040000

Epoch 16/30

- 0s - loss: 0.4291 - acc: 0.8950 - val\_loss: 0.6342 - val\_acc: 0.8070  
- LR: 0.038462

Epoch 17/30

- 0s - loss: 0.4138 - acc: 0.9060 - val\_loss: 0.6201 - val\_acc: 0.8110  
- LR: 0.037037

Epoch 18/30

- 0s - loss: 0.3986 - acc: 0.9070 - val\_loss: 0.6111 - val\_acc: 0.8130  
- LR: 0.035714

Epoch 19/30

- 0s - loss: 0.3859 - acc: 0.9140 - val\_loss: 0.6063 - val\_acc: 0.8110  
- LR: 0.034483

Epoch 20/30

- 0s - loss: 0.3740 - acc: 0.9140 - val\_loss: 0.5941 - val\_acc: 0.8130  
- LR: 0.033333

Epoch 21/30

- 0s - loss: 0.3638 - acc: 0.9170 - val\_loss: 0.5847 - val\_acc: 0.8210  
- LR: 0.032258  
Epoch 22/30  
- 0s - loss: 0.3536 - acc: 0.9220 - val\_loss: 0.5796 - val\_acc: 0.8200  
- LR: 0.031250  
Epoch 23/30  
- 0s - loss: 0.3440 - acc: 0.9270 - val\_loss: 0.5730 - val\_acc: 0.8280  
- LR: 0.030303  
Epoch 24/30  
- 0s - loss: 0.3361 - acc: 0.9250 - val\_loss: 0.5663 - val\_acc: 0.8250  
- LR: 0.029412  
Epoch 25/30  
- 0s - loss: 0.3282 - acc: 0.9280 - val\_loss: 0.5633 - val\_acc: 0.8280  
- LR: 0.028571  
Epoch 26/30  
- 0s - loss: 0.3213 - acc: 0.9280 - val\_loss: 0.5594 - val\_acc: 0.8270  
- LR: 0.027778  
Epoch 27/30  
- 0s - loss: 0.3145 - acc: 0.9290 - val\_loss: 0.5534 - val\_acc: 0.8280  
- LR: 0.027027  
Epoch 28/30  
- 0s - loss: 0.3077 - acc: 0.9340 - val\_loss: 0.5493 - val\_acc: 0.8310  
- LR: 0.026316  
Epoch 29/30  
- 0s - loss: 0.3022 - acc: 0.9350 - val\_loss: 0.5458 - val\_acc: 0.8280  
- LR: 0.025641  
Epoch 30/30  
- 0s - loss: 0.2966 - acc: 0.9350 - val\_loss: 0.5451 - val\_acc: 0.8260  
- LR: 0.025000

Baseline Error: 17.40%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])

As we can see the accuracy is around 83 percent on test data, 94 on training data and error is around 17 percent.

These are good metrics considering we used only 1000 images for our implementation, and it is well known that for neural networks at least we need millions of images to get a good model.

## References

<https://github.com/rajivranjanbuff>

## Question 2 (b)

Soln:

Below are the five files for this question:

two\_hidden\_layer.py--- two\_hidden\_layer without the L2 regularization

three\_hidden\_layer.py---three hidden layer without the L2 regularization  
single\_layer\_l2.py --- single layer with L2 regularization, lambda=5  
two\_hidden\_l2.py---two hidden layer with L2 regularization, lambda=5  
three\_hidden\_l2.py ---- three hidden layer with L2 regularization, lambda=5

output of two\_hidden\_layer.py--- two\_hidden\_layer without the L2 regularization :

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 30)	23550
dense_6 (Dense)	(None, 30)	930
dense_7 (Dense)	(None, 10)	310

Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

/Users/rajivranjan/anaconda3/lib/python3.6/site-packages/keras/models.py:942: UserWarning:  
The `nb\_epoch` argument in `fit` has been renamed `epochs`.

warnings.warn("The `nb\_epoch` argument in `fit` "

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 0s - loss: 2.3276 - acc: 0.0860 - val\_loss: 2.3050 - val\_acc: 0.1000  
- LR: 0.099900

Epoch 2/30

- 0s - loss: 2.3274 - acc: 0.0860 - val\_loss: 2.3023 - val\_acc: 0.1000  
- LR: 0.099800

Epoch 3/30

- 0s - loss: 2.3234 - acc: 0.0760 - val\_loss: 2.3074 - val\_acc: 0.1000  
- LR: 0.099701

Epoch 4/30

- 0s - loss: 2.3163 - acc: 0.1050 - val\_loss: 2.3121 - val\_acc: 0.1000  
- LR: 0.099602

Epoch 5/30

- 0s - loss: 2.3197 - acc: 0.0740 - val\_loss: 2.3150 - val\_acc: 0.1000  
- LR: 0.099502

Epoch 6/30



- 0s - loss: 2.3120 - acc: 0.1020 - val\_loss: 2.3157 - val\_acc: 0.1000  
- LR: 0.099404

Epoch 7/30

- 0s - loss: 2.3113 - acc: 0.1010 - val\_loss: 2.2981 - val\_acc: 0.2000  
- LR: 0.099305

Epoch 8/30

- 0s - loss: 2.2979 - acc: 0.1200 - val\_loss: 2.3027 - val\_acc: 0.1000  
- LR: 0.099206

Epoch 9/30

- 0s - loss: 2.2922 - acc: 0.1230 - val\_loss: 2.2810 - val\_acc: 0.1160  
- LR: 0.099108

Epoch 10/30

- 0s - loss: 2.2695 - acc: 0.1590 - val\_loss: 2.2510 - val\_acc: 0.3130  
- LR: 0.099010

Epoch 11/30

- 0s - loss: 2.2224 - acc: 0.2380 - val\_loss: 2.1829 - val\_acc: 0.2680  
- LR: 0.098912

Epoch 12/30

- 0s - loss: 2.1192 - acc: 0.2810 - val\_loss: 2.0821 - val\_acc: 0.2720  
- LR: 0.098814

Epoch 13/30

- 0s - loss: 1.9850 - acc: 0.3000 - val\_loss: 1.9313 - val\_acc: 0.3510  
- LR: 0.098717

Epoch 14/30

- 0s - loss: 1.8203 - acc: 0.3570 - val\_loss: 1.7833 - val\_acc: 0.4370  
- LR: 0.098619

Epoch 15/30

- 0s - loss: 1.6430 - acc: 0.4190 - val\_loss: 1.6186 - val\_acc: 0.4290  
- LR: 0.098522

Epoch 16/30

- 0s - loss: 1.4836 - acc: 0.4680 - val\_loss: 1.4931 - val\_acc: 0.4740  
- LR: 0.098425

Epoch 17/30

- 0s - loss: 1.3652 - acc: 0.5280 - val\_loss: 1.4035 - val\_acc: 0.5040  
- LR: 0.098328

Epoch 18/30

- 0s - loss: 1.2758 - acc: 0.5620 - val\_loss: 1.3373 - val\_acc: 0.5580  
- LR: 0.098232

Epoch 19/30

- 0s - loss: 1.1999 - acc: 0.6110 - val\_loss: 1.2843 - val\_acc: 0.5510  
- LR: 0.098135

Epoch 20/30

- 0s - loss: 1.1305 - acc: 0.6290 - val\_loss: 1.2458 - val\_acc: 0.5630  
- LR: 0.098039

Epoch 21/30

- 0s - loss: 1.0697 - acc: 0.6780 - val\_loss: 1.1888 - val\_acc: 0.6120  
- LR: 0.097943

Epoch 22/30

- 0s - loss: 1.0111 - acc: 0.6900 - val\_loss: 1.1489 - val\_acc: 0.6200  
- LR: 0.097847

Epoch 23/30

- 0s - loss: 0.9568 - acc: 0.7320 - val\_loss: 1.0907 - val\_acc: 0.6600  
- LR: 0.097752

Epoch 24/30

- 0s - loss: 0.8968 - acc: 0.7480 - val\_loss: 1.0723 - val\_acc: 0.6430  
- LR: 0.097656

Epoch 25/30

- 0s - loss: 0.8435 - acc: 0.7680 - val\_loss: 1.0196 - val\_acc: 0.6770  
- LR: 0.097561

Epoch 26/30

- 0s - loss: 0.7967 - acc: 0.7780 - val\_loss: 0.9871 - val\_acc: 0.6970  
- LR: 0.097466

Epoch 27/30

- 0s - loss: 0.7526 - acc: 0.7990 - val\_loss: 0.9554 - val\_acc: 0.7060  
- LR: 0.097371

Epoch 28/30

- 0s - loss: 0.7097 - acc: 0.8140 - val\_loss: 0.9253 - val\_acc: 0.7130  
- LR: 0.097276

Epoch 29/30

- 0s - loss: 0.6722 - acc: 0.8160 - val\_loss: 0.9065 - val\_acc: 0.7220  
- LR: 0.097182



Epoch 30/30



- 0s - loss: 0.6397 - acc: 0.8230 - val\_loss: 0.8820 - val\_acc: 0.7210  
- LR: 0.097087

Baseline Error: 27.90%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])







output of three\_hidden\_layer.py--- three\_hidden\_layer without the L2 regularization:

Layer (type)	Output Shape	Param #
=====		

dense_36 (Dense)	(None, 30)	23550
------------------	------------	-------

---

dense_37 (Dense)	(None, 30)	930
------------------	------------	-----

---

dense_38 (Dense)	(None, 30)	930
------------------	------------	-----

---

dense_39 (Dense)	(None, 10)	310
------------------	------------	-----

=====

Total params: 25,720

Trainable params: 25,720

Non-trainable params: 0

---

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 1s - loss: 2.3134 - acc: 0.0830 - val\_loss: 2.3027 - val\_acc: 0.1000

- LR: 0.009091

Epoch 2/30

- 0s - loss: 2.3053 - acc: 0.0720 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.004762

Epoch 3/30

- 0s - loss: 2.3041 - acc: 0.0760 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.003226

Epoch 4/30

- 0s - loss: 2.3037 - acc: 0.0930 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.002439

Epoch 5/30

- 0s - loss: 2.3035 - acc: 0.0840 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.001961

Epoch 6/30

- 0s - loss: 2.3033 - acc: 0.0860 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.001639

Epoch 7/30

- 0s - loss: 2.3032 - acc: 0.0880 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.001408

Epoch 8/30

- 0s - loss: 2.3031 - acc: 0.0850 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.001235

Epoch 9/30

- 0s - loss: 2.3030 - acc: 0.0790 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.001099

Epoch 10/30

- 0s - loss: 2.3030 - acc: 0.0760 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.000990

Epoch 11/30

- 0s - loss: 2.3029 - acc: 0.0830 - val\_loss: 2.3026 - val\_acc: 0.1000

- LR: 0.000901

Epoch 12/30

- 0s - loss: 2.3029 - acc: 0.0870 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000826

Epoch 13/30

- 0s - loss: 2.3029 - acc: 0.0780 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000763

Epoch 14/30

- 0s - loss: 2.3029 - acc: 0.0830 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000709

Epoch 15/30

- 0s - loss: 2.3028 - acc: 0.0970 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000662

Epoch 16/30

- 0s - loss: 2.3028 - acc: 0.0980 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000621

Epoch 17/30

- 0s - loss: 2.3028 - acc: 0.1000 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000585

Epoch 18/30

- 0s - loss: 2.3028 - acc: 0.0900 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000552

Epoch 19/30

- 0s - loss: 2.3028 - acc: 0.0810 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000524

Epoch 20/30

- 0s - loss: 2.3028 - acc: 0.0840 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000498

Epoch 21/30

- 0s - loss: 2.3028 - acc: 0.0720 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000474

Epoch 22/30

- 0s - loss: 2.3028 - acc: 0.0890 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000452

Epoch 23/30

- 0s - loss: 2.3027 - acc: 0.0940 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000433

Epoch 24/30

- 0s - loss: 2.3027 - acc: 0.0870 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000415

Epoch 25/30

- 0s - loss: 2.3027 - acc: 0.0760 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000398

Epoch 26/30

- 0s - loss: 2.3027 - acc: 0.0920 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000383

Epoch 27/30

- 0s - loss: 2.3027 - acc: 0.0950 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000369

Epoch 28/30

- 0s - loss: 2.3027 - acc: 0.0890 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000356

Epoch 29/30


- 0s - loss: 2.3027 - acc: 0.0980 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000344

Epoch 30/30


- 0s - loss: 2.3027 - acc: 0.0920 - val\_loss: 2.3026 - val\_acc: 0.1000  
- LR: 0.000332

Baseline Error: 10.00%


dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])








learning speed of the hidden layer



Train Error



single\_layer\_l2.py --- single layer with L2 regularization, lambda=5  
output:

Layer (type)	Output Shape	Param #
dense_40 (Dense)	(None, 30)	23550
dense_41 (Dense)	(None, 10)	310

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

---

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 1s - loss: 2.2759 - acc: 0.1830 - val\_loss: 2.1882 - val\_acc: 0.4040  
- LR: 0.090909

Epoch 2/30

- 0s - loss: 2.0348 - acc: 0.4160 - val\_loss: 1.8895 - val\_acc: 0.5900  
- LR: 0.083333

Epoch 3/30

- 0s - loss: 1.6393 - acc: 0.6300 - val\_loss: 1.5381 - val\_acc: 0.5910  
- LR: 0.076923

Epoch 4/30

- 0s - loss: 1.3030 - acc: 0.7150 - val\_loss: 1.2791 - val\_acc: 0.6580  
- LR: 0.071429

Epoch 5/30

- 0s - loss: 1.0709 - acc: 0.7660 - val\_loss: 1.1096 - val\_acc: 0.7110  
- LR: 0.066667

Epoch 6/30

- 0s - loss: 0.9075 - acc: 0.8050 - val\_loss: 1.0011 - val\_acc: 0.7350  
- LR: 0.062500

Epoch 7/30

- 0s - loss: 0.7993 - acc: 0.8270 - val\_loss: 0.9057 - val\_acc: 0.7530  
- LR: 0.058824

Epoch 8/30

- 0s - loss: 0.7153 - acc: 0.8480 - val\_loss: 0.8433 - val\_acc: 0.7610  
- LR: 0.055556

Epoch 9/30

- 0s - loss: 0.6547 - acc: 0.8530 - val\_loss: 0.8031 - val\_acc: 0.7790  
- LR: 0.052632

Epoch 10/30

- 0s - loss: 0.6037 - acc: 0.8640 - val\_loss: 0.7634 - val\_acc: 0.7810  
- LR: 0.050000

Epoch 11/30

- 0s - loss: 0.5657 - acc: 0.8730 - val\_loss: 0.7301 - val\_acc: 0.7930  
- LR: 0.047619

Epoch 12/30

- 0s - loss: 0.5325 - acc: 0.8760 - val\_loss: 0.7075 - val\_acc: 0.8000  
- LR: 0.045455

Epoch 13/30

- 0s - loss: 0.5047 - acc: 0.8840 - val\_loss: 0.6888 - val\_acc: 0.8000  
- LR: 0.043478

Epoch 14/30

- 0s - loss: 0.4806 - acc: 0.8890 - val\_loss: 0.6733 - val\_acc: 0.8010  
- LR: 0.041667

Epoch 15/30

- 0s - loss: 0.4597 - acc: 0.8970 - val\_loss: 0.6584 - val\_acc: 0.8110  
- LR: 0.040000

Epoch 16/30

- 0s - loss: 0.4413 - acc: 0.8950 - val\_loss: 0.6463 - val\_acc: 0.8080  
- LR: 0.038462

Epoch 17/30

- 0s - loss: 0.4262 - acc: 0.9060 - val\_loss: 0.6324 - val\_acc: 0.8130

- LR: 0.037037

Epoch 18/30

- 0s - loss: 0.4112 - acc: 0.9060 - val\_loss: 0.6235 - val\_acc: 0.8130

- LR: 0.035714

Epoch 19/30

- 0s - loss: 0.3987 - acc: 0.9130 - val\_loss: 0.6189 - val\_acc: 0.8110

- LR: 0.034483

Epoch 20/30

- 0s - loss: 0.3869 - acc: 0.9150 - val\_loss: 0.6067 - val\_acc: 0.8130

- LR: 0.033333

Epoch 21/30

- 0s - loss: 0.3769 - acc: 0.9170 - val\_loss: 0.5975 - val\_acc: 0.8210

- LR: 0.032258

Epoch 22/30

- 0s - loss: 0.3668 - acc: 0.9210 - val\_loss: 0.5924 - val\_acc: 0.8200

- LR: 0.031250

Epoch 23/30

- 0s - loss: 0.3574 - acc: 0.9270 - val\_loss: 0.5860 - val\_acc: 0.8270

- LR: 0.030303

Epoch 24/30

- 0s - loss: 0.3496 - acc: 0.9250 - val\_loss: 0.5794 - val\_acc: 0.8260

- LR: 0.029412

Epoch 25/30

- 0s - loss: 0.3418 - acc: 0.9280 - val\_loss: 0.5764 - val\_acc: 0.8270

- LR: 0.028571

Epoch 26/30

- 0s - loss: 0.3350 - acc: 0.9280 - val\_loss: 0.5727 - val\_acc: 0.8270

- LR: 0.027778

Epoch 27/30

- 0s - loss: 0.3283 - acc: 0.9290 - val\_loss: 0.5667 - val\_acc: 0.8290

- LR: 0.027027

Epoch 28/30

- 0s - loss: 0.3216 - acc: 0.9340 - val\_loss: 0.5628 - val\_acc: 0.8310

- LR: 0.026316

Epoch 29/30


- 0s - loss: 0.3162 - acc: 0.9350 - val\_loss: 0.5594 - val\_acc: 0.8280  
- LR: 0.025641



Epoch 30/30



- 0s - loss: 0.3107 - acc: 0.9350 - val\_loss: 0.5588 - val\_acc: 0.8260  
- LR: 0.025000

Baseline Error: 17.40%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])







two\_hidden\_l2.py---two hidden layer with L2 regularization, lambda=5

Layer (type)	Output Shape	Param #
=====		



dense_42 (Dense)	(None, 30)	23550
------------------	------------	-------

---

dense_43 (Dense)	(None, 30)	930
------------------	------------	-----

---

dense_44 (Dense)	(None, 10)	310
------------------	------------	-----

=====

Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

---

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 1s - loss: 2.3337 - acc: 0.0860 - val\_loss: 2.3111 - val\_acc: 0.1000

- LR: 0.099900

Epoch 2/30

- 0s - loss: 2.3334 - acc: 0.0860 - val\_loss: 2.3083 - val\_acc: 0.1000

- LR: 0.099800

Epoch 3/30

- 0s - loss: 2.3294 - acc: 0.0760 - val\_loss: 2.3134 - val\_acc: 0.1000

- LR: 0.099701

Epoch 4/30

- 0s - loss: 2.3223 - acc: 0.1050 - val\_loss: 2.3180 - val\_acc: 0.1000

- LR: 0.099602

Epoch 5/30

- 0s - loss: 2.3257 - acc: 0.0740 - val\_loss: 2.3209 - val\_acc: 0.1000

- LR: 0.099502

Epoch 6/30

- 0s - loss: 2.3181 - acc: 0.1010 - val\_loss: 2.3217 - val\_acc: 0.1000

- LR: 0.099404

Epoch 7/30

- 0s - loss: 2.3176 - acc: 0.1010 - val\_loss: 2.3045 - val\_acc: 0.2000

- LR: 0.099305

Epoch 8/30

- 0s - loss: 2.3046 - acc: 0.1190 - val\_loss: 2.3095 - val\_acc: 0.1000

- LR: 0.099206

Epoch 9/30

- 0s - loss: 2.2997 - acc: 0.1180 - val\_loss: 2.2889 - val\_acc: 0.1100

- LR: 0.099108

Epoch 10/30

- 0s - loss: 2.2789 - acc: 0.1570 - val\_loss: 2.2616 - val\_acc: 0.3160  
- LR: 0.099010

Epoch 11/30

- 0s - loss: 2.2363 - acc: 0.2250 - val\_loss: 2.1993 - val\_acc: 0.2680  
- LR: 0.098912

Epoch 12/30

- 0s - loss: 2.1403 - acc: 0.2780 - val\_loss: 2.1048 - val\_acc: 0.2720  
- LR: 0.098814

Epoch 13/30

- 0s - loss: 2.0113 - acc: 0.2960 - val\_loss: 1.9582 - val\_acc: 0.3480  
- LR: 0.098717

Epoch 14/30

- 0s - loss: 1.8519 - acc: 0.3470 - val\_loss: 1.8153 - val\_acc: 0.4220  
- LR: 0.098619

Epoch 15/30

- 0s - loss: 1.6786 - acc: 0.4090 - val\_loss: 1.6510 - val\_acc: 0.4230  
- LR: 0.098522

Epoch 16/30

- 0s - loss: 1.5159 - acc: 0.4600 - val\_loss: 1.5206 - val\_acc: 0.4700  
- LR: 0.098425

Epoch 17/30

- 0s - loss: 1.3933 - acc: 0.5200 - val\_loss: 1.4278 - val\_acc: 0.5000  
- LR: 0.098328

Epoch 18/30

- 0s - loss: 1.3023 - acc: 0.5550 - val\_loss: 1.3611 - val\_acc: 0.5510  
- LR: 0.098232

Epoch 19/30

- 0s - loss: 1.2265 - acc: 0.5990 - val\_loss: 1.3087 - val\_acc: 0.5460  
- LR: 0.098135

Epoch 20/30

- 0s - loss: 1.1578 - acc: 0.6180 - val\_loss: 1.2706 - val\_acc: 0.5480  
- LR: 0.098039

Epoch 21/30

- 0s - loss: 1.0983 - acc: 0.6640 - val\_loss: 1.2153 - val\_acc: 0.6070  
- LR: 0.097943

Epoch 22/30

- 0s - loss: 1.0413 - acc: 0.6840 - val\_loss: 1.1766 - val\_acc: 0.6120  
- LR: 0.097847

Epoch 23/30

- 0s - loss: 0.9885 - acc: 0.7230 - val\_loss: 1.1203 - val\_acc: 0.6540  
- LR: 0.097752

Epoch 24/30

- 0s - loss: 0.9293 - acc: 0.7460 - val\_loss: 1.1028 - val\_acc: 0.6380  
- LR: 0.097656

Epoch 25/30

- 0s - loss: 0.8762 - acc: 0.7600 - val\_loss: 1.0502 - val\_acc: 0.6750  
- LR: 0.097561

Epoch 26/30

- 0s - loss: 0.8294 - acc: 0.7700 - val\_loss: 1.0168 - val\_acc: 0.6910  
- LR: 0.097466

Epoch 27/30

- 0s - loss: 0.7853 - acc: 0.7970 - val\_loss: 0.9851 - val\_acc: 0.7060  
- LR: 0.097371

Epoch 28/30

- 0s - loss: 0.7427 - acc: 0.8100 - val\_loss: 0.9552 - val\_acc: 0.7090  
- LR: 0.097276

Epoch 29/30



- 0s - loss: 0.7059 - acc: 0.8140 - val\_loss: 0.9366 - val\_acc: 0.7140  
- LR: 0.097182



Epoch 30/30


- 0s - loss: 0.6746 - acc: 0.8190 - val\_loss: 0.9129 - val\_acc: 0.7160  
- LR: 0.097087

Baseline Error: 28.40%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])







output of three\_hidden\_l2.py ---- three hidden layer with L2 regularization, lambda=5

Layer (type)	Output Shape	Param #
dense_46 (Dense)	(None, 30)	23550
dense_47 (Dense)	(None, 30)	930
dense_48 (Dense)	(None, 30)	930
dense_49 (Dense)	(None, 10)	310

Total params: 25,720

Trainable params: 25,720

Non-trainable params: 0

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 1s - loss: 2.3351 - acc: 0.0830 - val\_loss: 2.3130 - val\_acc: 0.1000

- LR: 0.099010

Epoch 2/30

- 0s - loss: 2.3376 - acc: 0.0820 - val\_loss: 2.3154 - val\_acc: 0.1000

- LR: 0.098039

Epoch 3/30

- 0s - loss: 2.3357 - acc: 0.0790 - val\_loss: 2.3136 - val\_acc: 0.1000  
- LR: 0.097087

Epoch 4/30

- 0s - loss: 2.3282 - acc: 0.1010 - val\_loss: 2.3226 - val\_acc: 0.1000  
- LR: 0.096154

Epoch 5/30

- 0s - loss: 2.3324 - acc: 0.0940 - val\_loss: 2.3248 - val\_acc: 0.1000  
- LR: 0.095238

Epoch 6/30

- 0s - loss: 2.3249 - acc: 0.0990 - val\_loss: 2.3303 - val\_acc: 0.1000  
- LR: 0.094340

Epoch 7/30

- 0s - loss: 2.3307 - acc: 0.0990 - val\_loss: 2.3114 - val\_acc: 0.1000  
- LR: 0.093458

Epoch 8/30

- 0s - loss: 2.3277 - acc: 0.0800 - val\_loss: 2.3146 - val\_acc: 0.1000  
- LR: 0.092593

Epoch 9/30

- 0s - loss: 2.3229 - acc: 0.0880 - val\_loss: 2.3205 - val\_acc: 0.1000  
- LR: 0.091743

Epoch 10/30

- 0s - loss: 2.3272 - acc: 0.0820 - val\_loss: 2.3128 - val\_acc: 0.1000  
- LR: 0.090909

Epoch 11/30

- 0s - loss: 2.3247 - acc: 0.0780 - val\_loss: 2.3104 - val\_acc: 0.1000  
- LR: 0.090090

Epoch 12/30

- 0s - loss: 2.3230 - acc: 0.1000 - val\_loss: 2.3212 - val\_acc: 0.1000  
- LR: 0.089286

Epoch 13/30

- 0s - loss: 2.3268 - acc: 0.0810 - val\_loss: 2.3152 - val\_acc: 0.1000  
- LR: 0.088496

Epoch 14/30

- 0s - loss: 2.3223 - acc: 0.0930 - val\_loss: 2.3130 - val\_acc: 0.1000  
- LR: 0.087719

Epoch 15/30

- 0s - loss: 2.3205 - acc: 0.0910 - val\_loss: 2.3186 - val\_acc: 0.1000  
- LR: 0.086957

Epoch 16/30

- 0s - loss: 2.3255 - acc: 0.0880 - val\_loss: 2.3134 - val\_acc: 0.1000  
- LR: 0.086207

Epoch 17/30

- 0s - loss: 2.3211 - acc: 0.0960 - val\_loss: 2.3135 - val\_acc: 0.1000  
- LR: 0.085470

Epoch 18/30

- 0s - loss: 2.3220 - acc: 0.0840 - val\_loss: 2.3113 - val\_acc: 0.1000  
- LR: 0.084746

Epoch 19/30

- 0s - loss: 2.3224 - acc: 0.0850 - val\_loss: 2.3125 - val\_acc: 0.1000  
- LR: 0.084034

Epoch 20/30

- 0s - loss: 2.3218 - acc: 0.0860 - val\_loss: 2.3126 - val\_acc: 0.1000  
- LR: 0.083333

Epoch 21/30

- 0s - loss: 2.3226 - acc: 0.0780 - val\_loss: 2.3107 - val\_acc: 0.1000  
- LR: 0.082645

Epoch 22/30

- 0s - loss: 2.3174 - acc: 0.1070 - val\_loss: 2.3152 - val\_acc: 0.1000  
- LR: 0.081967

Epoch 23/30

- 0s - loss: 2.3196 - acc: 0.1080 - val\_loss: 2.3120 - val\_acc: 0.1000  
- LR: 0.081301

Epoch 24/30

- 0s - loss: 2.3153 - acc: 0.0970 - val\_loss: 2.3198 - val\_acc: 0.1000  
- LR: 0.080645

Epoch 25/30

- 0s - loss: 2.3225 - acc: 0.0930 - val\_loss: 2.3123 - val\_acc: 0.1000  
- LR: 0.080000



Epoch 26/30

- 0s - loss: 2.3181 - acc: 0.1140 - val\_loss: 2.3118 - val\_acc: 0.1000  
- LR: 0.079365

Epoch 27/30

- 0s - loss: 2.3184 - acc: 0.0880 - val\_loss: 2.3116 - val\_acc: 0.1000  
- LR: 0.078740

Epoch 28/30

- 0s - loss: 2.3185 - acc: 0.0920 - val\_loss: 2.3106 - val\_acc: 0.1000  
- LR: 0.078125

Epoch 29/30

- 0s - loss: 2.3185 - acc: 0.0900 - val\_loss: 2.3110 - val\_acc: 0.1000  
- LR: 0.077519


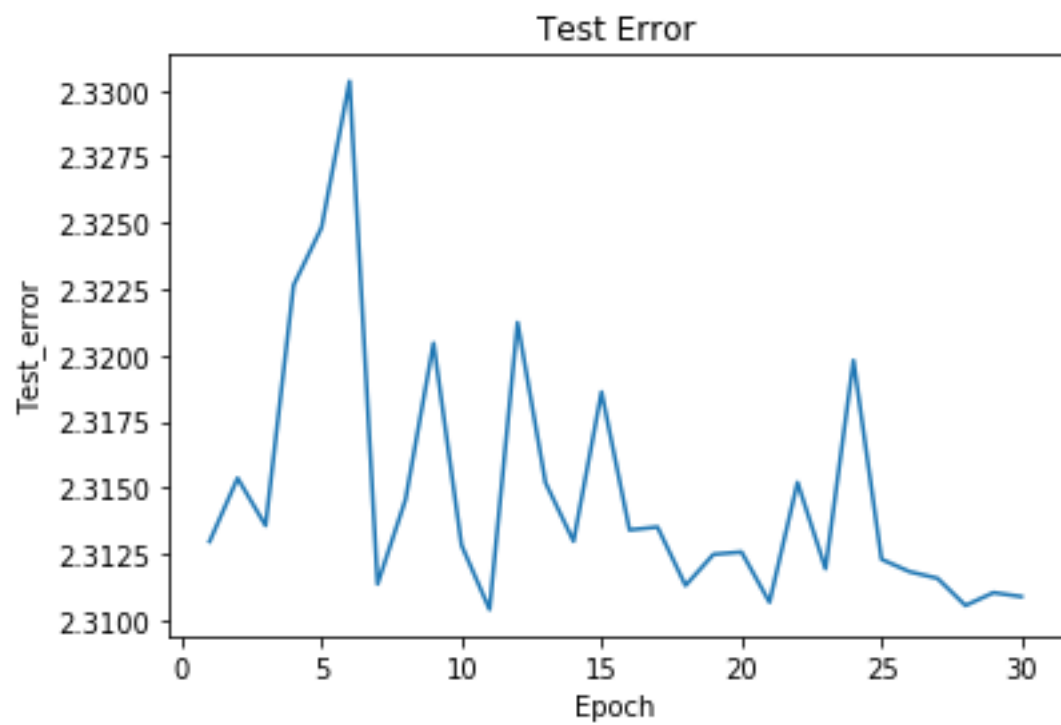
Epoch 30/30

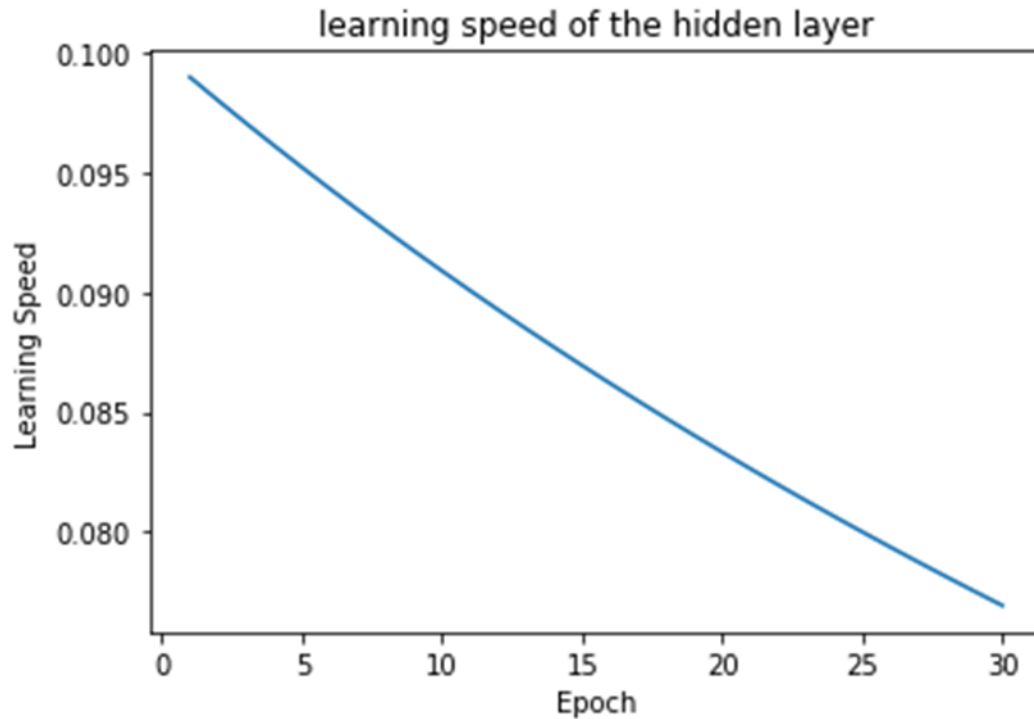
- 0s - loss: 2.3189 - acc: 0.0860 - val\_loss: 2.3109 - val\_acc: 0.1000  
- LR: 0.076923

Baseline Error: 10.00%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])







### Question 2 (c)

Construct and train convolutional neural network for MNIST classification. Regularize the training of the neural network through dropout. Regularize the training of neural network through augment your selection of 1000 images by rotating them for 1-3 degrees clockwise and counter clockwise, and shifting them for 3 pixels in 8 different directions. You can find many tutorials on those techniques, and our emphasize is that we understand those techniques.

Soln:

Construct and train convolutional neural network for MNIST classification.  
This is without dropout.

Filename: **MNIST\_convo.py**

This was run on the whole 60000 training data and the 10000 testing data

It was run for 10 epochs. The output is below:

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 171s - loss: 0.2331 - acc: 0.9334 - val\_loss: 0.0800 - val\_acc: 0.9756

Epoch 2/10

- 163s - loss: 0.0661 - acc: 0.9805 - val\_loss: 0.0640 - val\_acc: 0.9785

Epoch 3/10

- 162s - loss: 0.0452 - acc: 0.9864 - val\_loss: 0.0422 - val\_acc: 0.9856

Epoch 4/10

- 160s - loss: 0.0338 - acc: 0.9895 - val\_loss: 0.0395 - val\_acc: 0.9859

Epoch 5/10

- 163s - loss: 0.0273 - acc: 0.9913 - val\_loss: 0.0385 - val\_acc: 0.9869

Epoch 6/10

- 185s - loss: 0.0213 - acc: 0.9936 - val\_loss: 0.0376 - val\_acc: 0.9875

Epoch 7/10

- 234s - loss: 0.0163 - acc: 0.9954 - val\_loss: 0.0396 - val\_acc: 0.9874

Epoch 8/10

- 214s - loss: 0.0124 - acc: 0.9962 - val\_loss: 0.0377 - val\_acc: 0.9887

Epoch 9/10

- 212s - loss: 0.0097 - acc: 0.9972 - val\_loss: 0.0359 - val\_acc: 0.9887

Epoch 10/10

- 176s - loss: 0.0088 - acc: 0.9973 - val\_loss: 0.0369 - val\_acc: 0.9887

CNN Error: 1.13%

THE accuracy is 98.87 percent and the error is 1.13%.

With Dropout of 20 percent: File name is ***MNIST\_convo\_dropout.py***

This was run on the whole 60000 training data and the 10000 testing data

It was run for 10 epochs. The output is below:

Train on 60000 samples, validate on 10000 samples

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

- 223s - loss: 0.2315 - acc: 0.9343 - val\_loss: 0.0815 - val\_acc: 0.9743

Epoch 2/10

- 198s - loss: 0.0738 - acc: 0.9781 - val\_loss: 0.0469 - val\_acc: 0.9839

Epoch 3/10

- 178s - loss: 0.0532 - acc: 0.9839 - val\_loss: 0.0425 - val\_acc: 0.9862

Epoch 4/10

- 180s - loss: 0.0403 - acc: 0.9879 - val\_loss: 0.0402 - val\_acc: 0.9869

Epoch 5/10

- 186s - loss: 0.0336 - acc: 0.9894 - val\_loss: 0.0341 - val\_acc: 0.9883

Epoch 6/10

- 185s - loss: 0.0273 - acc: 0.9915 - val\_loss: 0.0301 - val\_acc: 0.9899

Epoch 7/10

- 171s - loss: 0.0233 - acc: 0.9927 - val\_loss: 0.0342 - val\_acc: 0.9886

Epoch 8/10

- 166s - loss: 0.0202 - acc: 0.9938 - val\_loss: 0.0324 - val\_acc: 0.9882

Epoch 9/10

- 167s - loss: 0.0169 - acc: 0.9944 - val\_loss: 0.0297 - val\_acc: 0.9901

Epoch 10/10

- 164s - loss: 0.0142 - acc: 0.9960 - val\_loss: 0.0316 - val\_acc: 0.9910

CNN Error: 0.90%

**As we see that with a dropout of 20 percent the accuracy increases:  
Now the accuracy is 99.10 percent and the error is just 0.90 percent.**

c) Regularize the training of neural network through augment your selection of 1000 images by rotating them for 1-3 degrees clockwise and counter clockwise, and shifting them for 3 pixels in 8 different directions. You can find many tutorials on those techniques, and our emphasize is that we understand those techniques.

Soln: The output is in file:**rotated\_neural.py**

④

(c) Regularize the training of neural network through augment your selection of 1000 images by rotating them for  $\pm 5$  degrees clockwise and counter clockwise and shifting them for 3 pixels in 8 different directions.

Soln:- In python, we can rotate an image by using the function

```
out = rotate(input_img, image degree, reshape=False)
```


$\rightarrow$  degree\_rot = 3 (counter-clockwise rotation of 3 degree)

degree\_rot = 357 ( $360 - 3$ )

$\rightarrow$  clockwise rotation

of 3 degree

shifting in 8 directions: for each of 1000 images



The shifting is as follows:-

0,0

0,28

28,0

28,28

left shift rows, 14, 15, 16

right shift <sup>shift in this direction</sup> rows 22, 23, 24  
(direction)

top shift  $\uparrow$  columns 17, 18, 19

bottom shift  $\downarrow$  columns 14, 15, 16

4 directions,

4 directions



For this firstly both the diagonals were found out & then they were rotated in both the directions.  $(2+2) = 4$  directions

Hence the image was counterclockwise rotated by 3 degree & 3 pixels shifted in 8 direction, right, left, top, bottom,

top right, top left, bottom right, bottom left. Hence, these were implemented



The output of the execution is below:

Layer (type)	Output Shape	Param #
dense_58 (Dense)	(None, 30)	23550
dense_59 (Dense)	(None, 10)	310

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

- 1s - loss: 2.2702 - acc: 0.1840 - val\_loss: 2.1824 - val\_acc: 0.4030  
- LR: 0.090909

Epoch 2/30

- 0s - loss: 2.0296 - acc: 0.4150 - val\_loss: 1.8840 - val\_acc: 0.5890  
- LR: 0.083333

Epoch 3/30

- 0s - loss: 1.6351 - acc: 0.6260 - val\_loss: 1.5323 - val\_acc: 0.5910  
- LR: 0.076923

Epoch 4/30

- 0s - loss: 1.2986 - acc: 0.7150 - val\_loss: 1.2722 - val\_acc: 0.6600  
- LR: 0.071429

Epoch 5/30

- 0s - loss: 1.0656 - acc: 0.7660 - val\_loss: 1.1014 - val\_acc: 0.7120  
- LR: 0.066667

Epoch 6/30

- 0s - loss: 0.9016 - acc: 0.8070 - val\_loss: 0.9920 - val\_acc: 0.7340  
- LR: 0.062500

Epoch 7/30

- 0s - loss: 0.7928 - acc: 0.8260 - val\_loss: 0.8961 - val\_acc: 0.7540  
- LR: 0.058824

Epoch 8/30

- 0s - loss: 0.7084 - acc: 0.8480 - val\_loss: 0.8332 - val\_acc: 0.7620  
- LR: 0.055556

Epoch 9/30

- 0s - loss: 0.6473 - acc: 0.8520 - val\_loss: 0.7925 - val\_acc: 0.7800  
- LR: 0.052632

Epoch 10/30

- 0s - loss: 0.5960 - acc: 0.8650 - val\_loss: 0.7525 - val\_acc: 0.7820  
- LR: 0.050000

Epoch 11/30

- 0s - loss: 0.5576 - acc: 0.8720 - val\_loss: 0.7189 - val\_acc: 0.7960  
- LR: 0.047619

Epoch 12/30

- 0s - loss: 0.5241 - acc: 0.8750 - val\_loss: 0.6959 - val\_acc: 0.8020  
- LR: 0.045455

Epoch 13/30

- 0s - loss: 0.4960 - acc: 0.8830 - val\_loss: 0.6770 - val\_acc: 0.7990  
- LR: 0.043478

Epoch 14/30

- 0s - loss: 0.4717 - acc: 0.8890 - val\_loss: 0.6613 - val\_acc: 0.8000  
- LR: 0.041667

Epoch 15/30

- 0s - loss: 0.4506 - acc: 0.8960 - val\_loss: 0.6461 - val\_acc: 0.8100  
- LR: 0.040000

Epoch 16/30

- 0s - loss: 0.4319 - acc: 0.8950 - val\_loss: 0.6338 - val\_acc: 0.8090  
- LR: 0.038462

Epoch 17/30

- 0s - loss: 0.4166 - acc: 0.9040 - val\_loss: 0.6197 - val\_acc: 0.8110  
- LR: 0.037037

Epoch 18/30

- 0s - loss: 0.4014 - acc: 0.9050 - val\_loss: 0.6106 - val\_acc: 0.8120  
- LR: 0.035714

Epoch 19/30

- 0s - loss: 0.3887 - acc: 0.9140 - val\_loss: 0.6057 - val\_acc: 0.8120  
- LR: 0.034483

Epoch 20/30

- 0s - loss: 0.3768 - acc: 0.9140 - val\_loss: 0.5936 - val\_acc: 0.8150  
- LR: 0.033333

Epoch 21/30

- 0s - loss: 0.3665 - acc: 0.9160 - val\_loss: 0.5843 - val\_acc: 0.8230  
- LR: 0.032258

Epoch 22/30

- 0s - loss: 0.3563 - acc: 0.9200 - val\_loss: 0.5790 - val\_acc: 0.8220  
- LR: 0.031250

Epoch 23/30

- 0s - loss: 0.3467 - acc: 0.9250 - val\_loss: 0.5726 - val\_acc: 0.8290  
- LR: 0.030303

Epoch 24/30

- 0s - loss: 0.3388 - acc: 0.9240 - val\_loss: 0.5658 - val\_acc: 0.8290  
- LR: 0.029412

Epoch 25/30

- 0s - loss: 0.3309 - acc: 0.9270 - val\_loss: 0.5629 - val\_acc: 0.8270  
- LR: 0.028571

Epoch 26/30

- 0s - loss: 0.3240 - acc: 0.9280 - val\_loss: 0.5590 - val\_acc: 0.8270  
- LR: 0.027778

Epoch 27/30

- 0s - loss: 0.3172 - acc: 0.9280 - val\_loss: 0.5529 - val\_acc: 0.8290  
- LR: 0.027027

Epoch 28/30

- 0s - loss: 0.3104 - acc: 0.9330 - val\_loss: 0.5489 - val\_acc: 0.8310  
- LR: 0.026316

Epoch 29/30

- 0s - loss: 0.3048 - acc: 0.9340 - val\_loss: 0.5454 - val\_acc: 0.8300  
- LR: 0.025641

Epoch 30/30

- 0s - loss: 0.2992 - acc: 0.9340 - val\_loss: 0.5447 - val\_acc: 0.8260  
- LR: 0.025000

Baseline Error: 17.40%

dict\_keys(['val\_loss', 'val\_acc', 'loss', 'acc'])

The accuracy is 82.60 percent and loss is 17.40 percent.